Folk Theorem and, in particular, at the thresholds where cooperation is more beneficial than defection in the game of a Prisoners' Dilemma.

Student: Sophie Shapcott

Supervisor: Dr Vincent Knight

Academic Year: 2019/20

School of Mathematics,
Cardiff University

A final year undergraduate project submitted in partial fulfiment of the requirements for MMORS (Masters in Mathematics, Operational Research and Statistics) taught programme.

SUMMARY

ACKNOWLEDGEMENTS

Contents

List of Figures

Chapter 1

Analyses

In this chapter, an analysis of the data collected via the methods described in the previous chapter (Chapter ??) is given. Firstly, a brief initial overview is provided, where descriptive statistics of the equilibria obtained and the overall characteristics of the strategies used is discussed. Following this, a critical analysis of the p-thresholds obtained is carried out. Here, the environmental effects, on the outcome of the game, discussed include: number and characteristics of opponents; noise; and degeneracy. Then a large-scale multivariate analysis is executed before considering the reliability of the collected data. Note, as of writing, the database currently has 825700 entries (rows) and a total number of 159 tournament sets.

1.1 Initial Analysis

In this section, all the data (including those games which could be degenerate) are considered. Taking a brief look at the graphs produced for each set of tournaments, it can be seen that the main 'shapes' obtained are as seen in Figure 1.1.

This will be discussed further in the next section as follows: Figures 1.1a and 1.1d, and general properties of the p-threshold, will be described in Section 1.2; the stochasticity and accuracy visible in Figure 1.1b will be analysed in Sections ?? and 1.4.2, respectively; and finally Figure 1.1c, and degeneracy overall, will be considered in Section 1.2.4.

The summary statistics gained from running the *describe* method of a pandas database is given in Figure ??. From this, it can be seen that the number of opponents the *Defector* played against ranged from one to seven, with an average

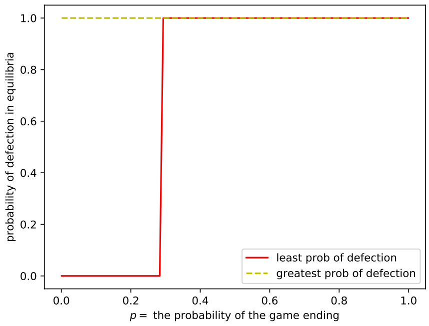

(a) An example of a graph with a clear p-threshold point of approximately 0.28. In this game there is no degeneracy and the opponent strategy playing in the tournament was *Inverse*, without noise.

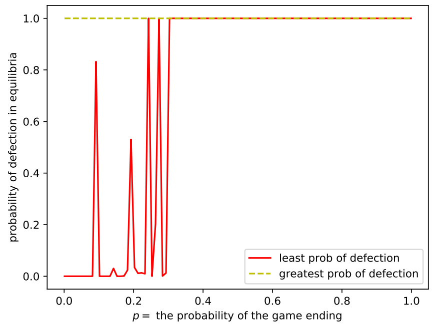

(b) An example of a graph where the p-threshold is not as clear. Perhaps this is due to a small amount of noise and hence not enough repetitions or stochasticity of the players. In this case the threshold seems to lie in the range [0.1, 0.3]. There is no degeneracy in this game and opponent strategies in the tournament were: *Feld: 1.0, 0.5, 200*; *Cooperator*; *EvolvedLookerUp2_2_2*; *Tullock: 11*; and *ZD-GEN-2: 0.125, 0.5, 3.* Again, this tournament was run with no added noise.

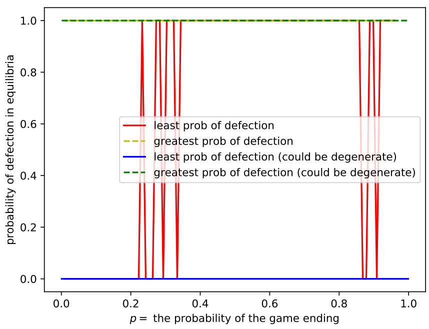

(c) An example of a graph where the p-threshold is not clear at all due to some (or all) of the games, which were played by the particular tournament set, being degenerate. In this case, the opponent strategies playing in the tournament were: *Random: 0.5*; *Grumpy: Nice, 10, -10*; *Fortress3*; and *Negation*. This tournament had no added noise.

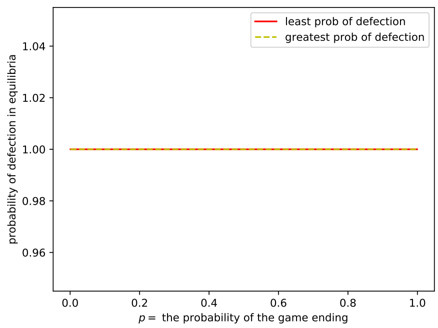

(d) An example of a graph where the least probability of defection was constant throughout the varying ending probabilities. In this particular tournament there was no ending probability in which cooperation was beneficial; with the least probability of defection equalling the greatest probability of defection at one. Here, the opponents were: *AntiCycler*; *e*; and *Stalker: (D,)*.

Figure 1.1: Example graphs obtained from the experiment.

of four opponents. Also, as expected, the mean probability of the game ending encountered was 0.5. Observe that, overall, there were $175,399$ distinct tournaments played (*experiment_number*) with a total of 159 distinct sets of strategies (*tournament_player_set*). Looking now at the statistics for Nash equilibria, it can be seen that a total of $823,823$ equilibrium points were calculated in this experiment, with an average of $1.914 \approx 2$ equilibria per game. However, observe, at least one game obtained 39 equilibria which will be explored into later on in this section. Considering the probabilities of defection within these equilibria, notice that both the greatest and the least probabilities of defection ranged from zero to one inclusive with a 50th percentile of zero. But, looking at the average values, the least probability has a mean of 0.342 and only just above this, the greatest probability has a mean of 0.460.

Next, further descriptive statistics are calculated for the strategies. This is to obtain a more in-depth view on the types of strategies randomly chosen to play and their characteristics. Executing *value_counts* method on the column of strategy names, it is observed that the player which appeared the most times (9 times) is *ZD-GEN-2: 0.125, 0.5, 3*; followed closely by *Tideman and Chieruzzi* with 7 sets of tournaments. On the other hand 38 out of the 200 strategies playing in this experiment appeared only once. Running the *value_counts* method again, but this time on the memory depths of the strategies found the majority of strategies to have an infinite memory depth. On the other hand, strategies having no memory or a depth equal to one were also significant. Considering the stochasticity of players alongside how many appearances each strategy made yielded the following chart in Figure 1.2.

It is clear that there is a clear bias towards deterministic strategies in this experiment. However, this is to be expected as running the following code:

```
len(axl.filtered_strategies(filterset={"stochastic": True})),
len(axl.filtered_strategies(filterset={"stochastic": False}))

(86, 156)
```

it can be seen that over half the strategies coded into the Axelrod library are classed as deterministic. Looking at Figure 1.2 again, observe, the majority of deterministic strategies were executed either once or three times. On the other hand, a large proportion of the stochastic strategies were played twice.

Further, the number of Nash equilibria obtained for each game was analysed and their distributions with respect to the number of opponents of the *Defector* plot-

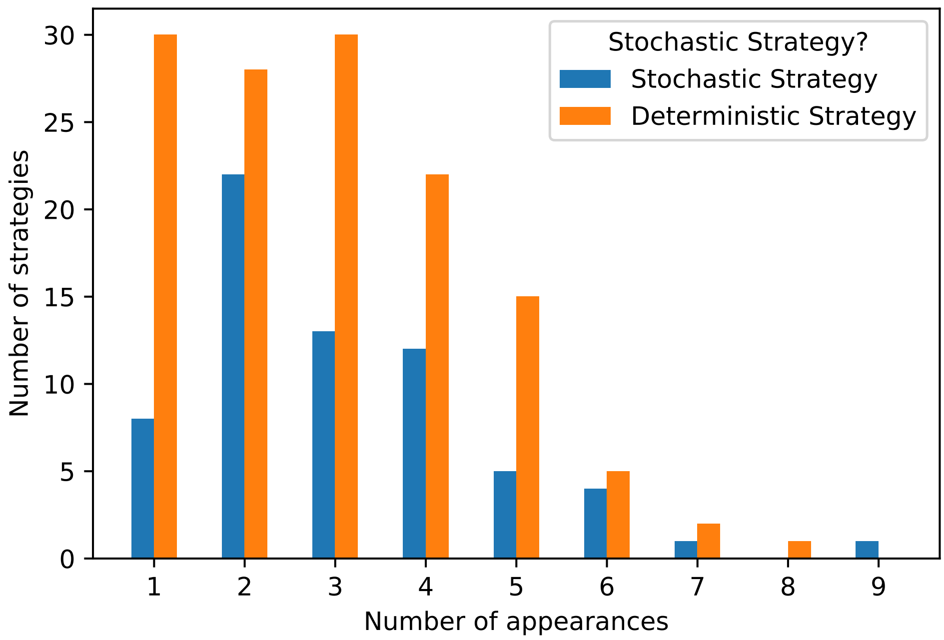

Figure 1.2: A plot to show the ratio of stochastic to deterministic strategies randomly chosen throughout the experiment.

ted. Executing the *value_counts* method on the 'num_of_equilibria' column gave the conclusion that the majority of games (131773) yielded one Nash equilibria with 28793 games obtaining 3 equilibria. Also, the maximum number of equilibria yielded, 39, was for one game, which, doing a search on the database, was found to be a six player game with noise=0.1. The opponents of the *Defector* were: *Inverse Punisher*; *Prober*; *PSO Gambler 2_2_2 Noise 05*; *Handshake*; and *More Tideman and Chieruzzi*.

Figure 1.3 shows the distributions of the number of Nash equilibria as per the number of players. As can be seen from the plot, the variance in the number of equilibria increases with the number of players, apart from when there were 6 players (5 opponents of the defector), where the spread is maximum. This could be due to the 39 equilibria gained for one game as discussed in the paragraph above. Looking now at the mean of the distributions, observe that these are also slightly increasing as the number of players increase.

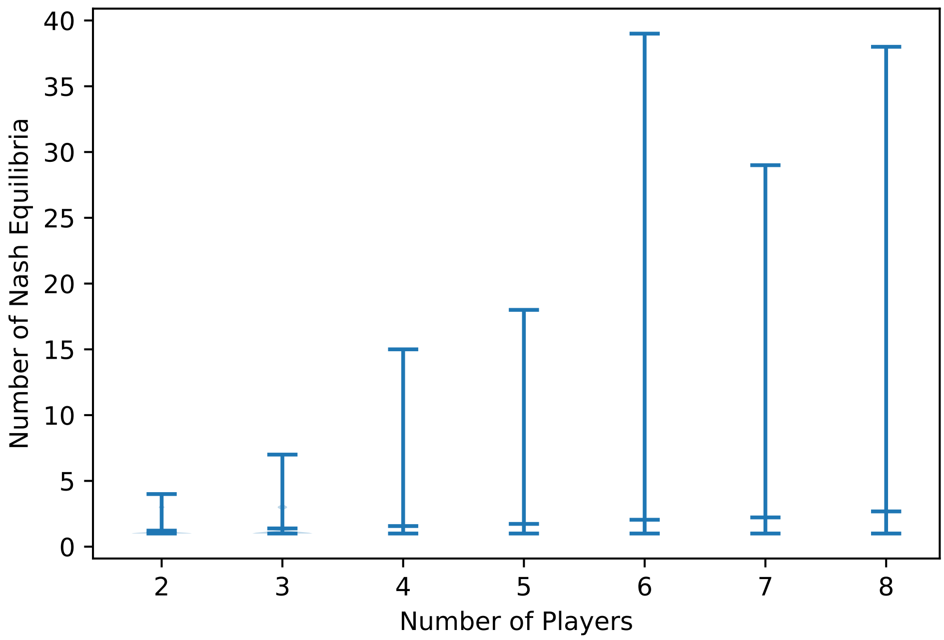

Figure 1.3: A violinplot showing the distribution of the number of equilibria obtained for varying number of players.

1.2 Analysis of the p-Threshold

In order to analyse the p-thresholds of the tournaments, two csv files were created [1] containing the minimum, mean, median and maximum probabilities for each set of tournaments. This was in order to gain as much information as possible from tournaments which gave graphs as shown in Figure 1.4. That is, tournaments in which the number of repetitions was not sufficient to omit the 'noise' which could affect the results.

For clarity, here is a restatement of the definition of a *p-threshold*: The maximum probability of the tournament ending for which cooperation is more beneficial than defection. That is the maximum probability of the tournament ending for which the probability of defecting is less than 0.5.

The first csv file created ignored whether a game could be degenerate and, by a small alteration, the second file contained only the non-degenerate games. This was in order to help identify whether degeneracy has any effect on the p-threshold. Within these files, other than the varying thresholds, the information about the

[1] Please see Appendix **??** for the code used to obtain these files. Also, those tournaments for which all the least probabilities of defection were greater than 0.5 were omitted - this indicated that regardless of the ending probability, it was more beneficial to defect.

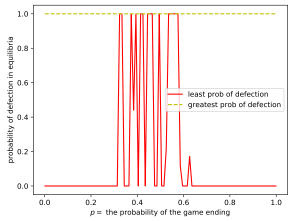

(a) This was an eight player tournament with no added noise. The opponent strategies were: *Tester*; *Gradual Killer: (D, D, D, D, D, C, C)*; *Win-Shift Lose-Stay: D*; *Evolved FSM 16*; *Retaliate: 0.1, 20*; *GEN-2: 0.125, 0.5, 3*; and *Joss: 0.9*.

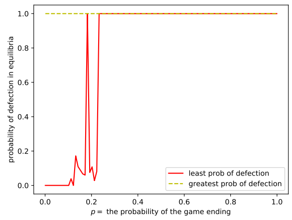

(b) This was a five player tournament with no added noise. The opponent strategies were: *WmAdams*; *Joss: 0.9*; *Limited Yamachi*; and *ZD- WSLS: 0.05*.

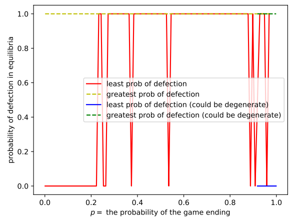

(c) This was a two player tournament with no added noise. The *Defector*'s opponent was *Stochastic Fortress3*.

Figure 1.4: Examples of plots where the p-threshold isn't as clear.

number of players, tournament strategy set and level of noise was retained.

Firstly, an exploration into the overall p-thresholds is given.

Observe, Figure 1.5a shows a 'split' in the data, with the majority of tournaments either having a threshold of less than 0.4 or a threshold equal to 1. This suggests that for a significant number of tournaments, it is more beneficial to cooperate if the probability of the tournament ending is sufficiently small enough (≤ 3). Moreover, there is a certain number of tournaments in which cooperation

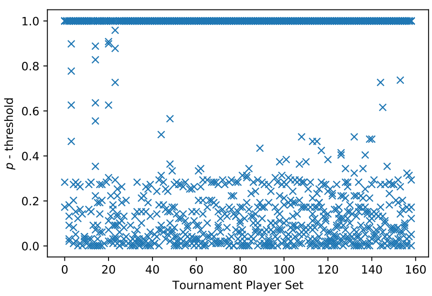

(a) A plot to show the minimum p-thresholds.

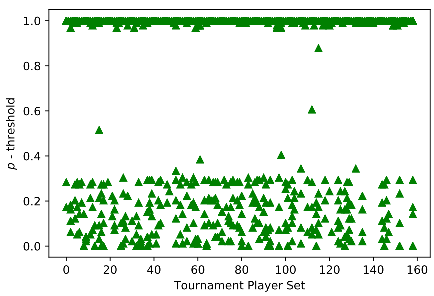

(b) A plot to show the maximum p-thresholds.

Figure 1.5: Plots to show the p-thresholds for all 1473 sets of tournaments.

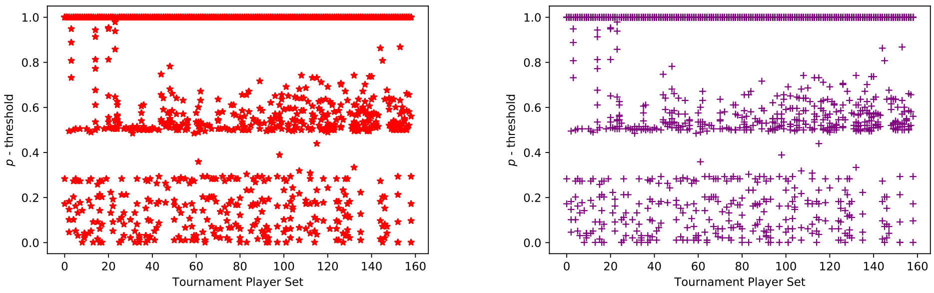

(a) A plot to show the mean p-thresholds. (b) A plot to show the median p-thresholds.

Figure 1.6: Plots to show the p-thresholds for all 1473 sets of tournaments.

is always preferential over defection. Similar inferences can be made from Figure 1.5b. On the other hand, looking at the plots in Figure 1.6, notice that there now appears a significant number of thresholds falling within the range $[0.5, 0.6]$. Looking at the tournaments which satisfy the latter observation (see Figure **??**), it can be seen that all of these tournaments had a wide variance. That is, a minimum threshold close to zero and a maximum threshold close to one. Indeed, looking closer at the plots created for a few of these tournaments, Figure 1.8, observe that there does not seem to be a clear p-threshold, with the plots having many peaks.

Consider now, the overall p-threshold for those tournaments which lead to a non-degenerate game. Comparing Figures 1.9a, 1.9b, 1.10a and 1.10b with Figures 1.5a, 1.5b, 1.6a and 1.6b, respectively, it can be seen that the majority of thresholds remain the same with only a very small number changing. This could be due to the number of games that were removed since obtaining the length of both csv files suggested that only four full tournaments were removed. The effects of degeneracy will be investigated in more depth in Section 1.2.4.

1.2.1 Effects of the Number of Players

In this section, the p-thresholds will be analysed with respect to the number of opponents the *Defector* played against. Note, in this section, only the mean p-thresholds of games which are non-degenerate will be discussed.

From Figure 1.11, it can be seen that the distributions of the p-thresholds with respect to the number of players are approximately the same. The means of the data falls around 0.7 with the main mode for each number of players being close to 1. Observe, though, that for the larger group of players (6, 7, 8), the distributions

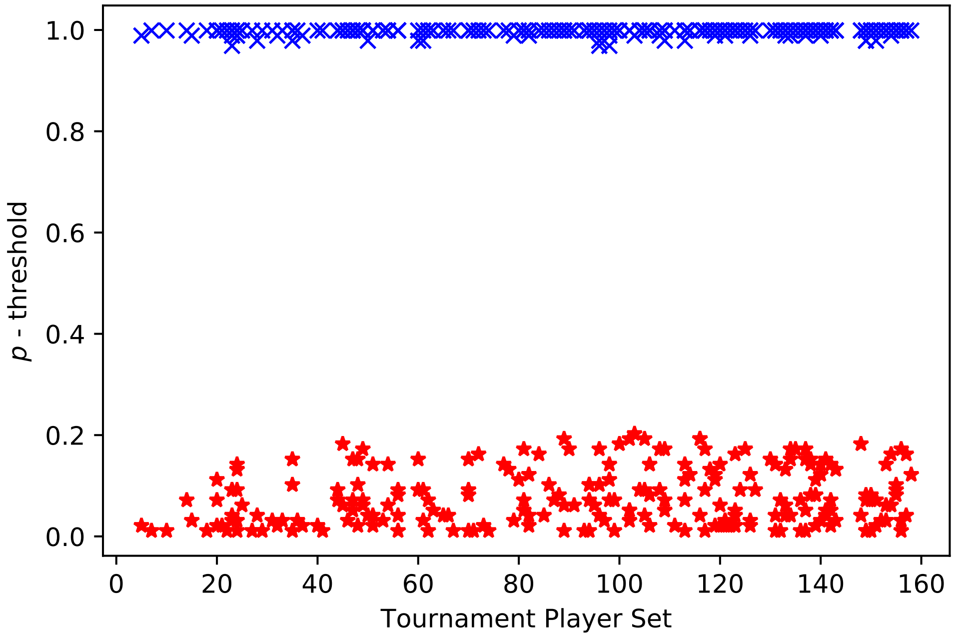

Figure 1.7: A plot to show the minimum and maximum p-thresholds for those tournaments which had an mean threshold within the range $[0.5, 0.6]$.

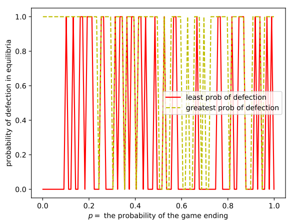

(a) This plot was obtained from the tournament in which the five opponents of the *Defector* were: *RichardHufford*; *Soft Joss: 0.9*; *Feld: 1.0, 0.5, 200*; *UsuallyDefects*; and *Handshake*. The tournament also included an additional 0.2 of noise.

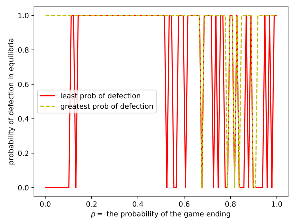

(b) This plot represents the, 0.1 noise, tournament with the three players: *Worse and Worse 2*; *Raider*; and, *Defector*.

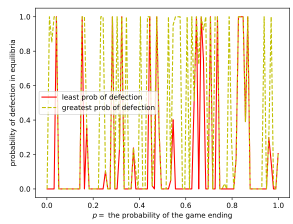

(c) This plot was produced from the tournament in which the *Defector* was playing *ZD-Extort-2: 0.1111111111111111, 0.5* and *Adaptive Tit For Tat: 0.5*. There was the addition of 0.5 noise within the tournament.

Figure 1.8: Example plots of the tournaments where the mean p-threshold was within the range $[0.5, 0.6]$.

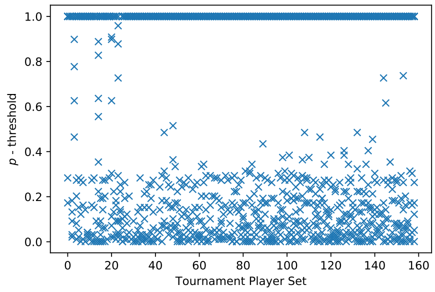

(a) A plot to show the minimum p-thresholds.

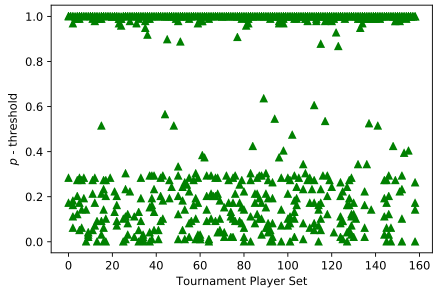

(b) A plot to show the maximum p-thresholds.

Figure 1.9: Plots to show the p-thresholds for all tournaments which led to non-degenerate games.

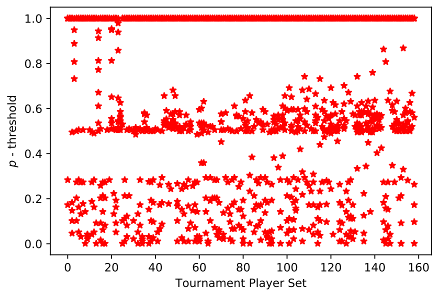

(a) A plot to show the mean p-thresholds.

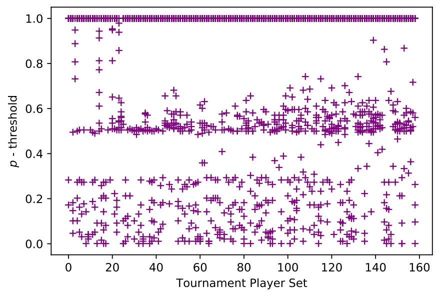

(b) A plot to show the median p-thresholds.

Figure 1.10: Plots to show the p-thresholds for all tournaments which led to non-degenerate games.

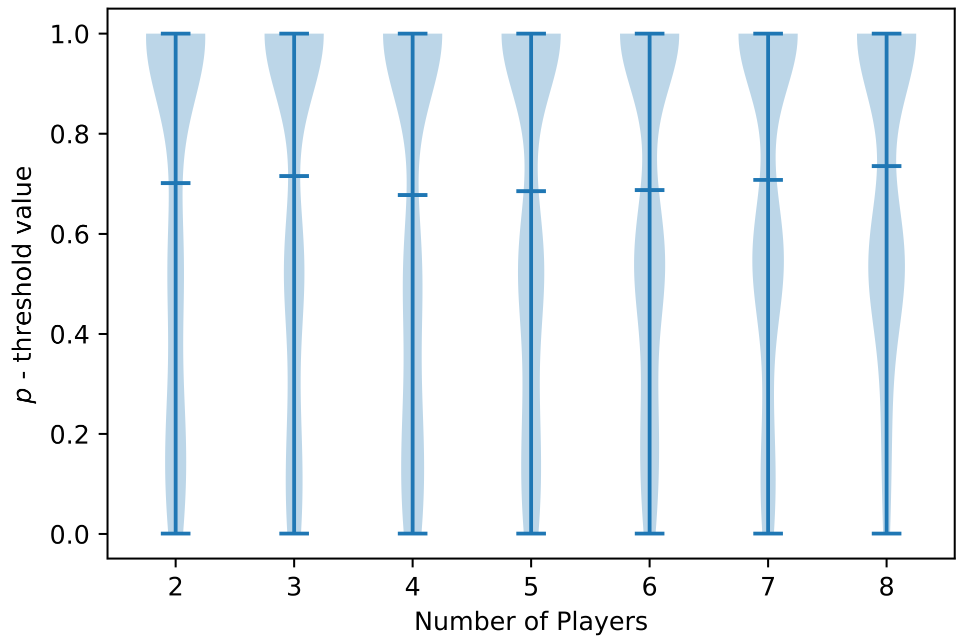

Figure 1.11: A violinplot showing the distribution of mean p-thresholds for each number of opponents.

become bimodal with the second mode around 0.5. This is most significant in the eight player tournaments however, note that, as of writing, the experiment was still running for eight player games and hence less data is available currently for this situation.

1.2.2 Effects of Stochastic Players

In order to analyse the effect of stochastic strategies on the p-threshold, tournaments in which these strategies featured are repeated but omitting the stochastic players to identify any alterations in the results.

1.2.3 Effects of Noise

Here, an analysis on the effects of noise on the p-threshold is provided. The addition of noise to a tournament indicates that, with a certain probability, the action of a particular strategy is altered [1]. That is, an action of *coop* changes to *defect* and vice versa.

Figure 1.12 shows the mean p-thresholds for all non-degenerate tournaments with the addition of varying probabilities of noise. Firstly, observe that for those tournaments with the addition of noise at least 0.6 (Figures 1.12g, 1.12g, 1.12g, 1.12g, 1.12g)

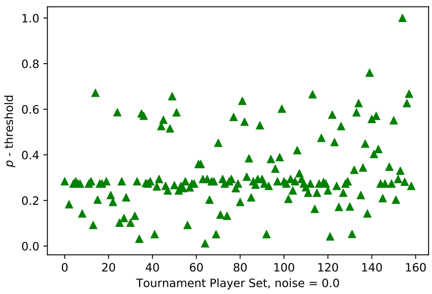

(a) A plot to show the mean p-thresholds for all tournaments with no added noise.

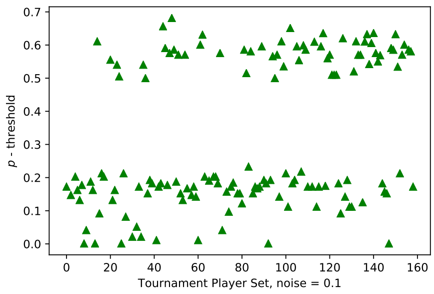

(b) A plot to show the mean p-thresholds for all tournaments with noise = 0.1.

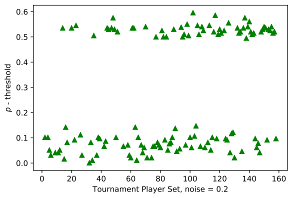

(c) A plot to show the mean p-thresholds for all tournaments with noise = 0.2.

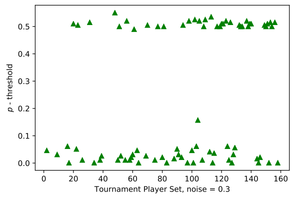

(d) A plot to show the mean p-thresholds for all tournaments with noise = 0.3.

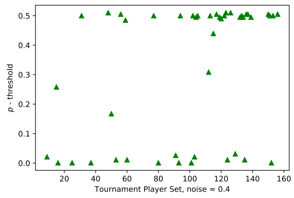

(e) A plot to show the mean p-thresholds for all tournaments with noise = 0.4.

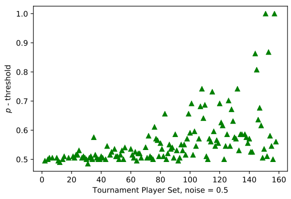

(f) A plot to show the mean p-thresholds for all tournaments with noise = 0.5.

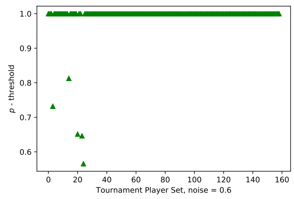

(g) A plot to show the mean p-thresholds for all tournaments with noise = 0.6.

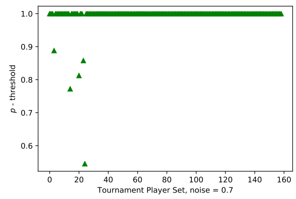

(h) A plot to show the mean p-thresholds for all tournaments with noise = 0.7.

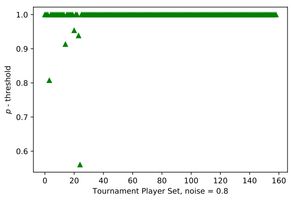

(i) A plot to show the mean p-thresholds for all tournaments with noise = 0.8.

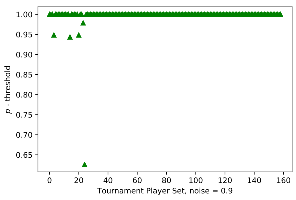

(j) A plot to show the mean p-thresholds for all tournaments with noise = 0.9.

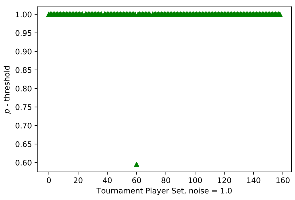

(k) A plot to show the mean p-thresholds for all tournaments with noise = 1.0.

Figure 1.12: Plots to show the mean p-threshold for varying levels of noise.

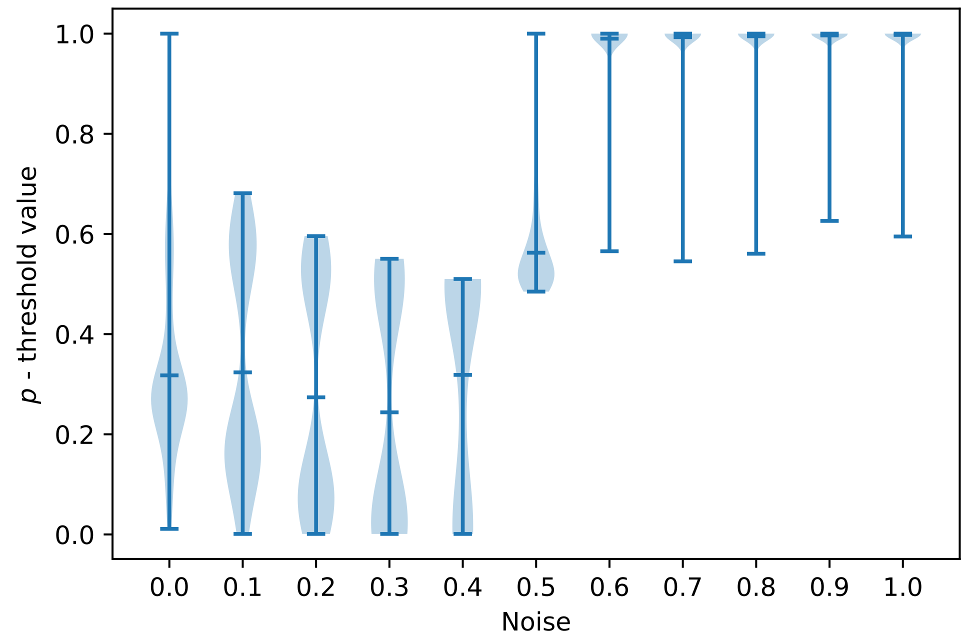

Figure 1.13: A violinplot showing the distribution of mean p-thresholds for each level of noise.

the majority of thresholds are equal to one. This is due to the definition of the parameter noise in the tournaments as described above, That is, there is a greater chance that the *Defector's* strategy will change to cooperation. Hence, only tournaments with noise$\leq$ 0.5 will be considered.

1.2.4 Effects of Degeneracy

1.3 Multivariate Data Analysis

1.4 Reliability of Data

1.4.1 Comparison of Databases

1.4.2 Accuracy of Thresholds

References

[1] Nikoleta E. Glynatsi and Vincent A. Knight. *A meta analysis of tournaments and an evaluation of performance in the Iterated Prisoner's Dilemma*. 2020. arXiv: `2001.05911 [cs.GT]`.